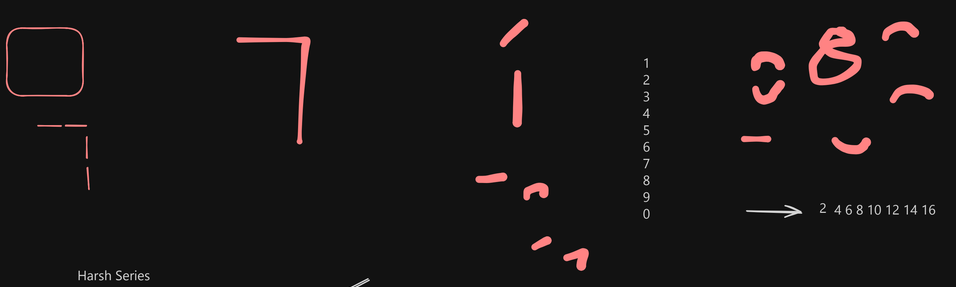





1
2
3
4
5
6
7
8
9
0



2 4 6 8 10 12 14 16

Harsh Series

Basic series

31 39 47 58 63 81 90

18 20 22 24 26 28 30

50 58 66 77 82 100 99

Harsh Series

Harsh Series

33 41 49 60 65 83 92

Hello How are you

Hello

How are you

wk wk wk
wq wq wq
wv wv wv

wk
wq
wv

Transformer

10*10

20*20

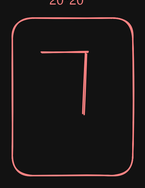

1*1

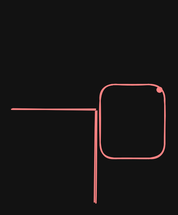

DSA | IQ| CGPA
10. 12. 8

60*60 = 3600 number

13 14 15 20 21 28 34

15 16 17 22 23 30 36

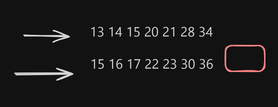

0-9

0-255

28*28 = 784 numbers

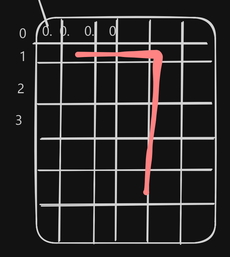

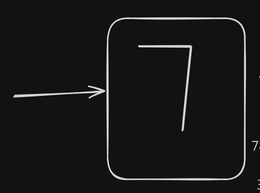

784 number

3600

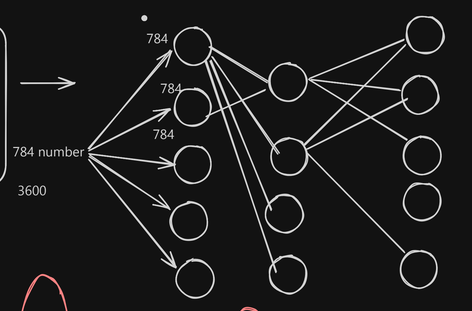

784
784
784

10 neuron

softmax: probability
Answer hoga

784 number

blacl

horizon

horizontal

3 horizontal
3 vertical

top curve
left curve
right curve
bottom curve

top curve
left curve
right curve
bottom curve

1 horizonta
1 vertical

circle
circle

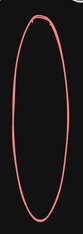

8

My name is Rohit ---

I am Rohit

Dog bites men ---?
men bites dog ---?

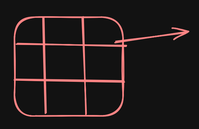

9 number

9 weights aapko learn karne hai

curve,
diagonal
horizontal
vertical
kuch nahi

9 number

9 weights aapko learn karne hai

0 5:
1 curve
2 horizontal
3 vertical
4 diagonal
blank

9 input

9 weight+B

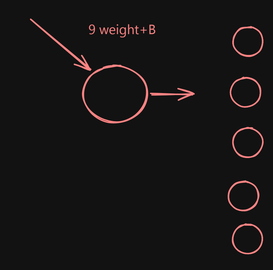

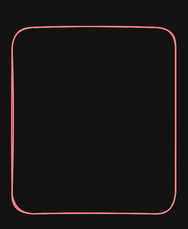

3*3 pixel: Every time

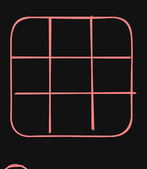

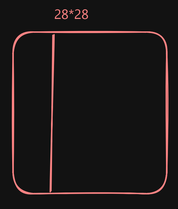

28*28

9 weight

9 weight

9 weight

9 weight

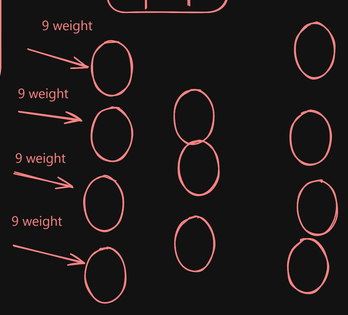

9 number

9 weights aapko learn karne hai

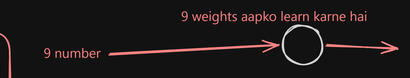

28*28 = 784

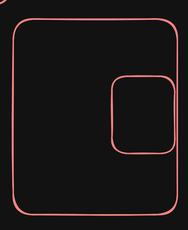

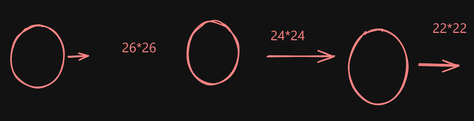

26*26

24*24

22*22

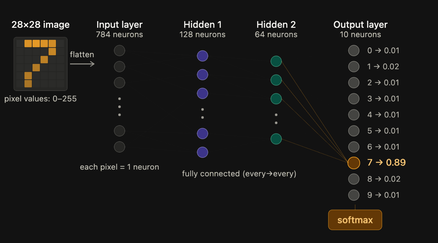


784 inputs × 128 = 100,352 weights in first layer alone
Every pixel connected to every neuron — no spatial awareness

The 28×28 image (that little grid on the left with a "7" drawn in amber pixels) gets flattened into a single column of 784 neurons. Row by row, pixel by pixel, all spatial information destroyed.
Then every single one of those 784 neurons connects to every neuron in Hidden Layer 1 (128 neurons). That's already 100,352 weights just in the first connection. Then 128×64 = 8,192 more weights. Then 64×10 = 640 weights to the output.

The output layer has 10 neurons — one per digit. Softmax converts the raw values into probabilities. The "7" neuron fires highest at 0.89.

Now — this works for 28×28. But think about what happens when the image is 224×224 (standard real-world size). That's 50,176 input neurons. First layer alone? 50,176 × 128 = 6.4 million weights. And it still has zero understanding of which pixels are neighbors.

We started with a 28×28 image. We applied 32 filters (each 3×3). Each filter slid across the entire image and produced a 26×26 feature map. So now we have 32 feature maps of 26×26.

Each feature map is like a new "image" — but instead of showing brightness, it shows "how strongly does my pattern match at each location." One map highlights horizontal edges, another highlights vertical edges, another highlights corners, etc.

Think about it — the 26×26 grid has 676 values per feature map. Across 32 maps, that's 21,632 values. And this is just from a tiny 28×28 image.

For a 224×224 image? Each feature map would be 222×222. With 32 filters, that's over 1.5 million values. Still a lot.

The inventor asked:

"Do I really need all 676 positions? If I detected a horizontal edge at pixel (10,10) and also at pixel (10,11) — those are basically the same detection. The edge didn't move much. Do I need both?"

The nearby values in a feature map are often very similar — redundant information. You could shrink the grid without losing important information.

How would you shrink it? You have a 26×26 grid of numbers and you want to make it smaller while keeping the important information. What's the simplest approach you can think of?

For example, if you have four neighboring values like:

0.8 0.6

0.3 0.2

28×28 image → convolution (32 filters) → 32 × 26×26 → max pooling → 32 × 13×13

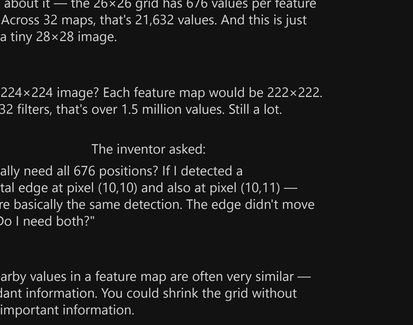


Conv1: 32 filters × (3×3×1) = 288
Conv2: 64 filters × (3×3×32) = 18,432
Conv3: 128 filters × (3×3×64) = 73,728
Conv4: 128 filters × (3×3×128) = 147,456